

# Viability-Preserving Rollout Approximate Dynamic Programming for Required-Target Orbital Inspection

**Rugang Tang**

School of Aeronautics, Northwestern Polytechnical University, Xi'an, Shaanxi 710072, China

Department of Aeronautical and Aviation Engineering, The Hong Kong Polytechnic University, Hung Hom, Kowloon, Hong Kong SAR, China

**Xin Ning**

School of Astronautics, Northwestern Polytechnical University, Xi'an, Shaanxi 710072, China

**Chih-Yung Wen**

Department of Aeronautical and Aviation Engineering, The Hong Kong Polytechnic University, Hung Hom, Kowloon, Hong Kong SAR, China

**Abstract**— Autonomous orbital inspection requires completing prescribed observations under relative-motion and safety constraints. This paper develops rollout approximate dynamic programming (ADP) on a directed graph of audited orbital arcs. The state records covered targets, selected viewpoints, and residual budget. A finite Bellman prefix uses the completion cost of a deterministic task-aware base policy as its approximate terminal value. Uncertified completions receive infinite value, while safe connecting actions remain available even without immediate inspection gain. Under deterministic finite-graph assumptions, exhaustive rollout preserves a completion certificate and does not exceed the cost of a completing base policy. Offline experiments use fixed geometry-defined requirements on an International Space Station mesh, with mandatory targets distinguished from optional background coverage. A nine-target confirmation retains 50 test

This work was supported in part by the Research Centre for Unmanned Autonomous Systems of The Hong Kong Polytechnic University under Grant P0046487. (Corresponding author: Rugang Tang).

Rugang Tang is with the School of Aeronautics, Northwestern Polytechnical University, Xi'an, Shaanxi 710072, China, and also with the Department of Aeronautical and Aviation Engineering, The Hong Kong Polytechnic University, Hung Hom, Kowloon, Hong Kong SAR, China (e-mail: [rugangtang@gmail.com](mailto:rugangtang@gmail.com)). Xin Ning is with the School of Astronautics, Northwestern Polytechnical University, Xi'an, Shaanxi 710072, China. Chih-Yung Wen is with the Department of Aeronautical and Aviation Engineering, The Hong Kong Polytechnic University, Hung Hom, Kowloon, Hong Kong SAR, China.

Supplementary material containing data and code is available with this article. Color versions of one or more figures are available online at <https://ieeexplore.ieee.org>.

0018-9251 © IEEE

and 30 shifted scenarios, including cases that lose required visibility. Depth-three ADP completes 43 and 9 missions, respectively. On 43 jointly completed test missions, mean maneuver velocity increment decreases by 11.05% relative to one-step-ADP-seeded local search. The paired reduction is 1.744 m/s, with a 95% interval of 1.525–1.923 m/s. A depth-one-to-six diagnostic quantifies the additional computation required for further cost improvement. The results support completion-certified ADP sequencing for prescribed inspection tasks.

**Index Terms**— spacecraft inspection, approximate dynamic programming, safe observable orbital arc, relative orbital dynamics, inspection requirements

## I. Introduction

Inspection of large orbital structures supports on-orbit servicing and proximity operations. External configuration changes, impact damage, and docking-interface anomalies may be difficult to observe with fixed onboard cameras [1]–[3]. An autonomous inspector must therefore decide where to observe and how to reach each observation. Relative dynamics couple these decisions because a useful camera pose may require an expensive or unsafe transfer.

Next-best-view methods model observation quality [4], kinodynamic planners combine coverage and differential constraints [1], [2], and relative-orbit planners exploit natural motion [5], [6]. A prescribed inspection mission also distinguishes mandatory surface targets from optional background coverage. An aggregate threshold can be reached while a required target remains unobserved. We therefore fix target identities and optimize the viewpoints and sequence needed to observe every prescribed item.

The planning challenge is to reduce motion cost while preserving a feasible continuation to all required observations. Each safe observable orbital arc (SOOA) combines a source-dependent Hill–Clohessy–Wiltshire (HCW) transfer, a stabilized camera observation, and feasibility audits. Safe connecting arcs remain useful even when they reveal no new target immediately. Fig. 1 presents the resulting workflow, from mission inputs and arc construction to rollout ADP and inspection records.

The main contribution is viability-preserving rollout approximate dynamic programming (ADP). A deterministic task-aware base policy supplies a completion cost, and finite-depth Bellman expansion improves the next action [7]. Replanning updates covered targets, selected observations, and residual budget after each action. Assigning infinite value to uncertified completions enforces the required task within this finite model. We establish completion preservation and base-policy no-degradation, with exact dynamic programming providing a reference on tractable graphs.

The contributions are:

- 1) a rollout ADP formulation that combines a required-target terminal predicate, finite-depth Bellman improvement, and a complete task-aware base-policy evaluation;




Fig. 1. Viability-preserving rollout ADP for orbital inspection. I. ISS mesh and required targets  $\mathcal{K}$  (orange rings). II. Camera-valid targets  $G_j$  from range, FOV, incidence  $\theta \leq \theta_{\max}$ , and LOS checks. III. HCW transfer, motion audits, and directed SOOA construction; C70 and C68 are archived candidate IDs. IV. Depth- $d$  Bellman lookahead with complete greedy base-policy tails. At each leaf,  $\sum \ell$  denotes the full tail cost;  $+\infty$  denotes not certified completion within the remaining budget. V. First-action selection and planning-state updates toward the required goal  $\Gamma_{\mathcal{K}}(m) = 1$ . The tree is schematic; the 9/9 route is an archived offline plan. ISS geometry: NASA. Camera photograph: A. Lucke, CC BY-SA 3.0 (illustrative hardware).

- 2) proofs of Markov sufficiency, preservation of a feasible task completion, base-policy no-degradation, and exact optimality on exhausted reduced graphs;
- 3) a directed SOOA representation with full-mesh visibility and motion audits, including complete GL Transmission Format (gITF) scene-node transforms, that admits safe connecting actions without requiring immediate inspection gain; and
- 4) a reproducible comparison on fixed required-target profiles, with matched baseline tasks, explicit coverage denominators, and a separate Robot Operating System 2 (ROS 2) execution evidence layer.

## II. Related Work

The Clohessy–Wiltshire equations provide a linear relative-motion model near a circular chief orbit [8], [9]. They permit controlled comparisons of observation sequencing under a common dynamical model. Hou et al. study relative orbits for visual inspection of the China Space Station [5]. Zhou et al. combine near-natural relative ellipses with formation control around a rotating station [10]. These studies optimize orbital or formation motion, complementing the discrete inspection sequence considered here.

Robotic next-best-view planning models sensor range, incidence, occlusion, reconstruction quality, and motion cost [4]. Rapidly exploring random trees and asymptotically optimal sampling methods incorporate dynamics and collision avoidance in continuous state space [11], [12]. Spacecraft extensions address kinodynamic coverage with one or several inspectors [1], [2]; astrodynamics-informed sampling exploits natural relative motion during transfer search [6]. OrbInspect represents each audited transfer-observation pair as a finite action and retains accumulated target identities in the decision state.

ADP and rollout approximate sequential decisions when exact Bellman recursion becomes expensive [7], [13]. Rollout evaluates a base-policy completion and applies truncated Bellman improvement. Deep reinforcement learning (DRL) uses learned value or policy approximations, while constrained policy optimization controls expected constraint costs [14]. Shielding instead removes actions that violate a specified safety model [15]. Building on rollout and shielding, the present model-based method combines audited actions with explicit goal-reaching completion values.

Recent work further integrates spacecraft geometry, motion, and sensing. Apodaca et al. study inspection trajectories and coverage–energy trade-offs [3], while

Zhou et al. develop surface-aware tomographic paths [16]. Ogri et al. address image-feature tracking, information gain, and dwell through switched-system logic [17]. These approaches motivate coordinated planning across observation geometry, transfer dynamics, and execution-level sensing.

OrbInspect addresses the discrete decision layer through a finite-completion viability test coupled to rollout ADP. Each SOOA links an observation to its dynamically audited transfer, so inspection credit depends on an admissible incoming motion. The resulting formulation extends established rollout machinery to prescribed inspection requirements on a directed action graph.

### III. System and Observation Model

The planner operates offline. It receives a known station mesh, a fixed nonempty required-target set, an initial relative state, camera parameters, a keep-out model, and actuation limits. It returns a time-indexed translational reference, a camera-boresight stream, selected SOOA records, accepted target identities, and summary metrics. Table I states the modeling assumptions and their implications for interpreting both evidence layers.

#### A. Relative Dynamics

The inspected structure is fixed in a local-vertical/local-horizontal (LVLH) frame attached to a nominal circular chief orbit. Its axes are radial, along-track, and cross-track, respectively. The chaser state is  $\mathbf{x} = [\mathbf{r}^\top, \mathbf{v}^\top]^\top \in \mathbb{R}^6$ , where  $\mathbf{r} = [r_x, r_y, r_z]^\top$  and  $\mathbf{v} = [v_x, v_y, v_z]^\top$ . With mean motion  $n$  and commanded acceleration  $\mathbf{u} = [a_x, a_y, a_z]^\top$ , the HCW model is

$$\begin{aligned} \dot{\mathbf{x}} &= A_{\text{HCW}}\mathbf{x} + B\mathbf{u}, \\ A_{\text{HCW}} &= \begin{bmatrix} 0 & I \\ K_n & C_n \end{bmatrix}, \quad B = \begin{bmatrix} 0 \\ I \end{bmatrix}, \\ K_n &= \text{diag}(3n^2, 0, -n^2), \quad C_n = \begin{bmatrix} 0 & 2n & 0 \\ -2n & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}. \end{aligned} \quad (1)$$

For the continuous HCW model, piecewise-constant acceleration over  $\Delta t$  has the exact zero-order-hold map

$$\begin{aligned} \mathbf{x}_{k+1} &= A_d \mathbf{x}_k + B_d \mathbf{u}_k, \\ A_d &= e^{A_{\text{HCW}} \Delta t}, \\ B_d &= \int_0^{\Delta t} e^{A_{\text{HCW}} \tau} B d\tau. \end{aligned} \quad (2)$$

The implementation constructs and propagates each boundary map with fixed-step fourth-order Runge–Kutta integration under the same piecewise-constant input model. The numerical map is  $\mathbf{x}_{k+1} = \tilde{A}_d \mathbf{x}_k + \tilde{B}_d \mathbf{u}_k$ , with (2) providing the exact continuous-model reference. Each rest-to-rest segment uses finite-duration thrust and is accepted when its terminal position and velocity errors satisfy the prescribed tolerances.

#### B. Visibility, Prescribed Observation, and Coverage

The structure is represented by surface samples  $\mathcal{T} = \{(p_i, n_i, a_i)\}_{i=1}^N$  with represented areas  $a_i > 0$ . Mission priority weights are separately defined as  $w_i = a_i \eta_i$ , where  $\eta_i > 0$  is a priority multiplier; they are not geometric areas. A fixed nonempty set  $\mathcal{K} \subseteq \mathcal{T}$  specifies the required inspection items before scenario evaluation. The camera attitude  $q_k$  defines unit boresight  $b(q_k)$ , and  $\mathcal{O}$  denotes the occluding mesh. Target  $i$  is visible when

$$\begin{aligned} r_{\min}^{\text{cam}} &\leq \|p_i - \mathbf{r}_k\| \leq r_{\max}^{\text{cam}}, \\ \arccos \frac{(p_i - \mathbf{r}_k)^\top b(q_k)}{\|p_i - \mathbf{r}_k\|} &\leq \alpha_{\max}, \\ \arccos \frac{(\mathbf{r}_k - p_i)^\top n_i}{\|\mathbf{r}_k - p_i\| \|n_i\|} &\leq \theta_{\max}, \\ \text{LOS}(\mathbf{r}_k, p_i, \mathcal{O}) &= 1. \end{aligned} \quad (3)$$

The four tests exclude targets that are too near or too far, outside the field of view, observed at an excessive incidence angle, or blocked by station geometry. Line of sight is queried against a deterministic bounding-volume hierarchy containing every ISS mesh triangle; intersections at the target endpoint are excluded, while interior segment intersections block the view. In a closed-loop execution model, dwell could be accumulated as  $\tau_{i,k+1} = \tau_{i,k} + \Delta t$  while target  $i$  remains visible and accepted when  $\tau_i \geq \tau_{\min}$ . The offline study instead treats every candidate as a stabilized terminal observation with its prescribed exposure completed a priori. Selecting node  $j$  therefore credits its fixed visibility mask  $G_j$ , so  $m_{k+1} = m_k \vee G_j$ . The sampled-area coverage represented by mask  $m_k$  is

$$C^{\text{area}}(m_k) = \frac{\sum_i a_i m_{i,k}}{\sum_i a_i}. \quad (4)$$

Candidate generation can make a subset of the sampled surface invisible from every candidate observation. For the nonempty set  $\mathcal{T}_{\text{insp}} = \cup_{j=1}^M G_j$ , we also report priority-weighted inspectable coverage:

$$C_k^{\text{insp}} = \frac{\sum_{i \in \mathcal{T}_{\text{insp}}} w_i m_{i,k}}{\sum_{i \in \mathcal{T}_{\text{insp}}} w_i}. \quad (5)$$

Equation (4) estimates credited area over the original surface samples, whereas (5) is a priority-weighted score conditional on the candidate library. Their denominators remain fixed after scenario-specific observation loss. Neither is the primary required-target completion metric, which is

$$\begin{aligned} C^{\mathcal{K}}(m) &= \frac{|\mathcal{K} \cap \text{supp}(m)|}{|\mathcal{K}|}, \\ \Gamma_{\mathcal{K}}(m) &= \mathbf{1}\{\mathcal{K} \subseteq \text{supp}(m)\}. \end{aligned} \quad (6)$$

Mission success requires  $\Gamma_{\mathcal{K}} = 1$ , independently of optional area coverage. One observation can credit several required items, and different viewpoints can satisfy the same item. Completion denotes acceptance of the prescribed samples under (3). Reported mission duration is transfer-only; observation exposure is assumed complete at each stabilized candidate.

TABLE I  
Modeling assumptions and the boundary of the reported evidence

| Assumption                     | Role in the baseline                                                                                                                                               | Consequence for interpretation                                                                                                         |
|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| Circular chief orbit           | Supports linear HCW relative translation in the LVLH frame.                                                                                                        | Quantitative results compare planners under the same local linear model, not a perturbed flight ephemeris.                             |
| Station fixed in LVLH          | Keeps mesh targets and the mesh safety geometry stationary during one planning run.                                                                                | Rotating appendages or station attitude motion require a body-frame extension.                                                         |
| Known mesh and state           | Makes visibility, occlusion, and transfer evaluation deterministic.                                                                                                | Navigation, map, and shape uncertainty are not represented by the present safety margin.                                               |
| Prescribed camera attitude     | Points the boresight toward each candidate observation target and exports the corresponding quaternion stream.                                                     | Camera slew time, reaction-wheel limits, and closed-loop attitude error are outside the optimization.                                  |
| Mesh-based sampled safety      | Computes distance to the full International Space Station (ISS) triangle mesh at every rollout sample and rejects sampled motion segments that intersect the mesh. | Positive sampled clearance and surface nonintersection are planning audits rather than a robust continuous-time collision certificate. |
| No illumination or plume model | Isolates view geometry and translational feasibility.                                                                                                              | Sun angle, shadow, contamination, and plume-impingement restrictions must be added for flight design.                                  |

### C. Dynamic Feasibility and SOOAs

Every sampled state must satisfy

$$\|\mathbf{u}_k\|_2 \leq u_{\max}, \quad \|\mathbf{v}_k\|_2 \leq v_{\max}, \quad d_S(\mathbf{r}_k) \geq 0, \quad (7)$$

where  $d_S(\mathbf{r}) = d_{\text{mesh}}(\mathbf{r}) - r_{\text{veh}} - d_{\text{safe}}$  is finite-body clearance above the required margin,  $d_{\text{mesh}}$  is unsigned center distance to the closest triangle in the complete ISS mesh, and  $r_{\text{veh}} = 0.80 \text{ m}$  conservatively bounds the chaser visual model; thus  $d_S = 0$  is the sampled admissibility boundary. The mesh coordinates are obtained from the complete glTF scene hierarchy (a matrix or translation–rotation–scale at every ancestor), followed by the Simulation Description Format (SDF) pose and scale. In addition, each consecutive pair of rollout samples is tested for a triangle-surface intersection, and its continuous center-distance lower bound subtracts half the segment length from the smaller endpoint distance by the 1-Lipschitz property of distance to a closed set. This swept test catches a direct mesh crossing even if both endpoints have positive unsigned distance. For a directed rollout from source node  $i$  to destination node  $j$ , define  $\Delta v_{ij} = \sum_k \|\mathbf{u}_{ij,k}\|_2 \Delta t$ , root-mean-square distance  $e_{ij}^{\text{rms}}$  from the transfer samples to the destination position, and  $d_{ij}^{\min} = \min_k d_S(\mathbf{r}_{ij,k})$ . The graph stage cost used in the reported study is

$$\ell_{ij} = \lambda_v \Delta v_{ij} + \lambda_e e_{ij}^{\text{rms}} + \lambda_s \max(0, -d_{ij}^{\min}) + \lambda_a. \quad (8)$$

For every shield-admissible edge, the clearance-violation term is zero. It remains in the shared cost implementation so deliberately unshielded comparison methods can be evaluated with the same expression.

A candidate observation node and a directed SOOA edge are, respectively,

$$\begin{aligned} v_j &= (\bar{\mathbf{r}}_j, \bar{q}_j, G_j), \\ a_{ij} &= (\mathbf{x}_{ij}[1:K], \mathbf{u}_{ij}[0:K-1], q_{ij}[1:K], \\ &\quad \Delta v_{ij}, d_{ij}^{\min}, u_{ij}^{\max}, v_{ij}^{\max}, e_{ij}^{\text{rms}}, e_{ij}^v, \rho_{ij}). \end{aligned} \quad (9)$$

Here  $\bar{\mathbf{r}}_j$  and  $\bar{q}_j$  are the terminal observation pose and prescribed attitude,  $G_j \subseteq \mathcal{T}$  is its camera-valid target

set, and  $e_{ij}^{\text{rms}}$  and  $e_{ij}^v$  are terminal position and velocity errors. The stored state and attitude arrays contain the  $K$  post-step samples; the source state is held by node  $v_i$ . The initial condition is a distinguished source node  $v_0$  with no visibility reward. For passive horizon  $T_p = L_p \Delta t$ , an optional audit can propagate zero-input HCW drift from every stored edge state and record  $\rho_{ij} = \min_{1 \leq k \leq K, 0 \leq l \leq L_p} d_S(\mathbf{r}_{ij,k}^{\text{drift}}(l))$ . The reported study disables this optional audit and therefore makes no passive-safety claim; it retains input, speed, full-mesh sampled clearance, swept surface-intersection, and terminal checks. Throughout this paper, “safe” means admissible under the audits explicitly enabled for the reported experiment. These deterministic checks do not certify robustness between samples or under model and state uncertainty.

## IV. Viability-Preserving Rollout ADP

### A. Candidate Graph and Markov State

Let  $\mathcal{V} = \{v_0, v_1, \dots, v_M\}$  contain the initial node and the  $M$  candidate observation nodes in (9). A directed edge from  $v_i$  to  $v_j$  is the source-dependent SOOA  $a_{ij}$ , and it is absent when an enabled audit fails. Thus visibility  $G_j$  is node dependent, whereas trajectory, cost, clearance, input, terminal error, and passive margin are edge dependent. This separation is necessary because reaching the same observation from two different source nodes generally produces different HCW motion and audit values.

A node identifier is not a sufficient decision state because the value of a view changes after its targets have been covered. The finite-horizon Markov state is

$$s_k = (j_k, m_k, b_k, h_k), \quad (10)$$

where  $j_k \in \{0, \dots, M\}$  is the current node,  $m_k \in \{0, 1\}^N$  is the covered-target mask,  $b_k \in \{0, 1\}^M$  is the selected-observation mask, and  $h_k$  is the remaining action budget. Selecting destination  $a \in \{1, \dots, M\}$  traverses edge  $a_{j_k a}$

and gives the deterministic transition

$$f(s_k, a) = (a, m_k \vee G_a, b_k \vee \mathbf{e}_a, h_k - 1). \quad (11)$$

The required set  $\mathcal{K}$  is fixed problem data. The shield-admissible action set  $\mathcal{U}_s(s)$  is defined below. The constrained graph problem minimizes total stage cost subject to (6), with zero terminal cost for a completed task and value  $+\infty$  when no admissible completion exists within the remaining budget. Its Bellman recursion is

$$V^*(s) = \begin{cases} 0, & \Gamma_{\mathcal{K}}(m) = 1, \\ +\infty, & h = 0 \text{ or } \mathcal{U}_s(s) = \emptyset, \\ \min_{a \in \mathcal{U}_s(s)} Q^*(s, a), & \text{otherwise,} \end{cases} \quad (12)$$

$$Q^*(s, a) = \ell_{ja} + V^*(f(s, a)).$$

Here  $j$  and  $h$  are the current-node and remaining-budget components of  $s$ . Exact evaluation is exponential in the coverage and selected-node masks. The proposed method approximates this recursion with a short exact prefix and a computationally inexpensive, goal-aware terminal policy.

## B. SOOA Construction and Safety Shield

Area-weighted targets are sampled from the mesh, and camera seeds are generated on offset shells around their outward normals. The binary visibility matrix  $V_{ij} = \mathbf{1}\{i \in G_j\}$  is computed once from the terminal observation geometry. For a transition to  $\bar{\mathbf{x}}_j = [\mathbf{r}_j^T, \mathbf{0}^T]^T$ , the planner solves the minimum-energy piecewise-constant HCW boundary-value problem under the numerical fourth-order Runge–Kutta (RK4) map

$$\begin{aligned} \min_{\mathbf{u}_0, \dots, \mathbf{u}_{K-1}} \quad & \sum_{k=0}^{K-1} \|\mathbf{u}_k\|_2^2 \\ \text{s.t.} \quad & \mathbf{x}_{k+1} = \tilde{A}_d \mathbf{x}_k + \tilde{B}_d \mathbf{u}_k, \\ & \mathbf{x}_K = \bar{\mathbf{x}}_j. \end{aligned} \quad (13)$$

If the terminal influence matrix  $\mathcal{B}_K = [\tilde{A}_d^{K-1} \tilde{B}_d \dots \tilde{B}_d]$  has full row rank, the stacked minimum-norm command from source state  $\mathbf{x}_i^{\text{src}}$  is  $\mathbf{U}_{ij}^* = \mathcal{B}_K^T (\mathcal{B}_K \mathcal{B}_K^T)^{-1} (\bar{\mathbf{x}}_j - \tilde{A}_d^K \mathbf{x}_i^{\text{src}})$ . The implementation computes this expression through the  $6 \times 6$  Gram system and then propagates the requested input without clipping. Let  $\chi_{ij} = 1$  exactly when every stored edge sample satisfies (7),  $e_{ij}^r \leq \varepsilon_r$ ,  $e_{ij}^v \leq \varepsilon_v$ , and  $\rho_{ij} \geq 0$  when the passive audit is enabled. The shield then admits

$$\mathcal{U}_s(s) = \{a \in \{1, \dots, M\} : b_a = 0, \chi_{ja} = 1\}, \quad (14)$$

where  $j$  is the current-node component of  $s$ . Safe unvisited nodes remain admissible when they provide no immediate inspection gain, because a connecting action may be needed to reach a required target on the directed graph. The finite budget and no-revisit mask bound such connections. The shield makes an inadmissible edge unavailable to incumbent construction, local search, and exact recursion [15]. Its guarantee remains conditional on the sampled numerical rollout and deterministic model used to compute  $\chi_{ij}$ .

## Algorithm 1 Viability-preserving rollout ADP

---

**Require:** Candidate graph, shield  $\mathcal{U}_s$ , required set  $\mathcal{K}$ , budget  $H$ , depth  $d$

- 1:  $s \leftarrow (0, \mathbf{0}, \mathbf{0}, H)$ ;  $\pi \leftarrow ()$
- 2: **while**  $\Gamma_{\mathcal{K}}(m) = 0$  and  $h > 0$  **do**
- 3:   Enumerate every  $a \in \mathcal{U}_s(s)$
- 4:   **for** each admissible  $a$  **do**
- 5:      $\hat{Q}_d(s, a) \leftarrow \ell_{ja} + \hat{V}_{d-1}(f(s, a))$
- 6:   **end for**
- 7:   Discard every action with  $\hat{Q}_d(s, a) = +\infty$
- 8:   **if** no finite action remains **then**
- 9:     **return** no certified completion
- 10:   **end if**
- 11:    $a^* \leftarrow \arg \min_a \hat{Q}_d(s, a)$  with deterministic ties
- 12:   Append  $a^*$  and its audit record to  $\pi$
- 13:    $s \leftarrow f(s, a^*)$
- 14: **end while**
- 15: **return**  $\pi$  if  $\Gamma_{\mathcal{K}}(m) = 1$ ; otherwise return no certified completion

---

## C. Adaptive Base Policy and Rollout Value

Let  $g_{\mathcal{K}}(s, a) = |\mathcal{K} \cap (G_a \setminus \text{supp}(m))|$  count newly accepted required targets, and let  $\epsilon = 0.05$  prevent division by a near-zero stage cost. For the required-target task, the deterministic adaptive base policy is

$$\mu(s) \in \arg \max_{a \in \mathcal{U}_s(s)} \left( \frac{g_{\mathcal{K}}(s, a)}{\max\{\epsilon, \ell_{ja}\}}, g_{\mathcal{K}}(s, a), -\ell_{ja} \right), \quad (15)$$

where the tuple is compared lexicographically, followed by deterministic node-identifier tie breaking. The policy recomputes gain after every accepted observation; zero-gain ties favor lower stage cost. Starting from state  $s$ , its completion value is

$$\hat{V}_0(s) = \begin{cases} \sum_{k=0}^{L-1} \ell_{jk} \mu(s_k), & \mu \text{ reaches } \Gamma_{\mathcal{K}} = 1 \text{ within } h, \\ +\infty, & \text{otherwise.} \end{cases} \quad (16)$$

Thus  $\hat{V}_0$  approximates the cost-to-go through a complete, shielded base-policy simulation.

For rollout depth  $d \geq 1$ , define the truncated Bellman value

$$\hat{V}_d(s) = \begin{cases} 0, & \Gamma_{\mathcal{K}}(m) = 1, \\ +\infty, & h = 0, \\ \min_{a \in \mathcal{U}_s(s)} \hat{Q}_d(s, a), & \text{otherwise,} \end{cases} \quad (17)$$

where  $\hat{Q}_d(s, a) = \ell_{ja} + \hat{V}_{d-1}(f(s, a))$ . The minimum over an empty action set is  $+\infty$ . The policy selects a finite-valued minimizing action, applies (11), and rebuilds the depth- $d$  tree at the successor state. State–depth memoization removes duplicate subproblems within each decision. The guarantees below require exhaustive prefix evaluation and the same deterministic Markov base policy at every leaf. Fig. 2 gives a worked decision on an illustrative graph, separating motion admissibility from completion certification and showing the continuation retained after replanning.




Fig. 2. Illustrative completion-certified rollout ADP, with initial state  $s = (0, 0, 0, 3)$ . I. Solid edges are admissible, edge labels are stage costs, and  $G_j$  lists visible targets. The dashed transfer is excluded. II. At  $d = 2$ , rollout selects zero-gain connector  $v_2$  for total cost 4, versus base-policy cost 8; action 3 has no certified completion. III. The successor  $s' = (2, 0, e_2, 2)$  retains a cost-3 completion. Counts indicate covered required targets;  $J_{\text{rollout}}$  is realized graph cost. The example illustrates Propositions 4 and 5. Assigned costs and audit flags are illustrative, not experimental HCW data; the primary study uses  $d = 3$ .

## D. Properties and Exact Oracle

**PROPOSITION 1** (Monotone submodularity of fixed-mask coverage) *For an observation set  $S \subseteq \{1, \dots, M\}$ , let  $F(S) = \sum_{i=1}^N w_i \mathbf{1}\{i \in \cup_{j \in S} G_j\}$ . If all  $w_i \geq 0$ , then  $F$  is normalized, monotone, and submodular.*

*Proof:*

Clearly  $F(\emptyset) = 0$ . If  $A \subseteq B$ , every target covered by  $A$  is covered by  $B$ , so nonnegative weights give  $F(A) \leq F(B)$ . For  $a \notin B$ , the marginal gain  $F(A \cup \{a\}) - F(A)$  is the total weight in  $G_a \setminus \cup_{j \in A} G_j$ . Since  $A \subseteq B$ , this set contains  $G_a \setminus \cup_{j \in B} G_j$ , proving diminishing returns. ■

**Remark 1** (Role of submodularity) Proposition 1 also holds for required-target gain by assigning unit weight to  $\mathcal{K}$  and zero elsewhere. It does not justify removing zero-gain connecting actions or supply a greedy approximation ratio for a directed inspection tour. Both SOOA cost and the admissibility of the next observation depend on the current node.

**PROPOSITION 2** (Markov sufficiency) *For a fixed deterministic candidate graph with fixed node masks, required set  $\mathcal{K}$ , source-dependent edge records, target weights,*

*revisit rule, and horizon, the state in (10) is Markov for the finite graph decision problem.*

*Proof:*

The current node  $j_k$  identifies the source of every possible next edge;  $m_k$  and fixed  $\mathcal{K}$  determine the remaining inspection requirements;  $b_k$  determines which observations remain selectable under the no-revisit rule; and  $h_k$  determines the remaining budget. By assumption, visibility masks and edge costs and audits are deterministic records indexed only by the current and destination nodes. Equations (11), (12), and (14) therefore depend on no earlier history once  $(j_k, m_k, b_k, h_k)$  is given. The proposition concerns binary prescribed-observation acceptance; partial dwell, multiple independent look directions, or time-dependent illumination would require additional state. ■

**PROPOSITION 3** (Compositional preservation of stored audits) *Consider a sequence  $(a_{j_{k-1}j_k})_{k=1}^L$  whose first source is  $v_0$ , whose idealized rest-to-rest endpoints match exactly, and whose every edge has  $\chi_{j_{k-1}j_k} = 1$ . Then every stored sample of the concatenated sequence satisfies the input, speed, and clearance inequalities, every edge endpoint satisfies its terminal conditions, and every stored passive-drift rollout has nonnegative margin when that audit is enabled.*

*Proof:*

By the definition of  $\chi_{ij}$ , all stated conditions hold on each accepted edge record. Exact endpoint matching identifies the terminal state of edge  $k$  with the source state of edge  $k + 1$ , so the records concatenate without an unmodeled jump. The sample set of the complete trajectory is the finite union of the per-edge sample sets; therefore every member of that union retains the corresponding per-edge audit result. ■

**Remark 2** (Meaning of audit preservation) The proposition is a finite-record composition result, not a continuous-time barrier certificate. It does not cover inter-sample clearance, endpoint perturbations beyond tolerance, closed-loop tracking error, or robustness to model and state uncertainty.

**PROPOSITION 4** (Rollout viability) *For the deterministic finite graph, if  $\widehat{V}_d(s) < +\infty$ , Algorithm 1 selects a shield-admissible action whose successor has a finite goal-reaching completion within the remaining budget. Replanning preserves this property until the goal is reached.*

*Proof:*

Finiteness in (17) provides a certificate consisting of shield-admissible prefix actions followed, unless the goal is reached earlier, by a base-policy completion. The selected first action is part of such a certificate, so its successor retains a goal-reaching sequence within the reduced budget. When the depth- $d$  tree is rebuilt, the certificate tail remains admissible; if another exact prefix action is required, the next action of the certified base completion supplies it. Determinism preserves the recorded masks,

costs, and audits. Induction over the decreasing action budget therefore proves goal attainment. ■

**PROPOSITION 5** (Base-policy no-degradation) *If the adaptive base policy reaches the goal from  $s$ , then  $\widehat{V}_d(s)$  is finite and  $\widehat{V}_d(s) \leq \widehat{V}_0(s)$  for every  $d \geq 1$ . For a fixed depth  $d$ , the realized graph cost of repeated replanning in Algorithm 1 is no greater than  $\widehat{V}_d(s)$  and therefore no greater than the base-policy cost.*

*Proof:*

The first  $d$  base-policy actions, truncated if the goal is reached earlier, followed by the remaining base completion form a feasible certificate in (17); hence  $\widehat{V}_d(s) \leq \widehat{V}_0(s)$ . Now take a minimizing depth- $d$  certificate at a replanning state and let  $s'$  be the successor of its first action  $a$ . The certificate tail, extended when necessary by the next base-policy action, is feasible in the rebuilt depth- $d$  tree at  $s'$ . Therefore  $\ell_{ja} + \widehat{V}_d(s') \leq \widehat{V}_d(s)$ . Summing this inequality over successive replans telescopes to a realized cost no greater than the initial  $\widehat{V}_d(s)$ . The result concerns the graph stage cost; reductions in physical  $\Delta v$  are evaluated separately. ■

For reduced-graph certification, exact recursion caches (12) for every reachable tuple  $(j, m, b, h)$  and enumerates every action in  $\mathcal{U}_s(s)$  rather than a local reversal neighborhood.

**THEOREM 1** (Optimality on an exhausted reduced SOOA graph) *For fixed finite nodes, deterministic visibility masks, deterministic audited edge costs, a required set  $\mathcal{K}$ , the no-revisit rule, and horizon  $H$ , exhaustive evaluation of (12) returns a minimum-cost task-completing shield-feasible sequence when one exists. If no such sequence exists, it returns value  $+\infty$  and reports infeasibility.*

*Proof:*

Proceed by induction on the remaining budget  $h$ . At  $h = 0$ , (12) returns zero exactly for a goal state and  $+\infty$  otherwise, which is correct. Assume the recursion is correct for all states with at most  $h - 1$  remaining actions. At a state with  $h$  actions, every feasible sequence begins with exactly one  $a \in \mathcal{U}_s(s)$  and then a feasible continuation from  $f(s, a)$ ; by the induction hypothesis, the recursion assigns that continuation its minimum cost or  $+\infty$  if none exists. Minimizing over every admissible first action therefore returns the minimum total cost, or  $+\infty$  when all continuations are infeasible. The claim is limited to the exhausted sampled graph and does not extend to omitted camera poses or continuous-space trajectories. ■

With maximum branching factor  $B$ , depth  $d$ , and at most  $H$  base-policy actions, direct decision evaluation costs  $O(HB^{d+1})$  in the worst case. Memoization reduces repeated evaluation when prefixes reach the same mask state. Fig. 3 and Table III quantify the offline cost-computation trade-off for  $d = 1, \dots, 6$ .

## V. Experimental Protocol

### A. Fixed Geometry and Required-Target Profiles

The study uses the corrected public ISS mesh distributed by the National Aeronautics and Space Administration (NASA) [18]. Its complete transformed geometry contains 247 525 triangles. Ninety equal-represented-area samples define the original surface denominator; 41 are observable from at least one of the 24 retained candidate observations. The directed archive contains 24 initial-to-observation and  $24 \times 23$  inter-observation records, of which 264 pass every enabled gate. Every method receives the same graph and scenario-specific available-node set.

Required-target profiles are selected deterministically from geometry before scenario evaluation. The mission manifest records their IDs, coordinates, selection rule, visibility masks, and checksums. These synthetic profiles distribute prescribed samples across the observable library; they are not an operational ISS maintenance checklist.

The required set remains fixed when a scenario removes observation nodes. Losing every available view of a required item makes the task impossible in the library. Otherwise, a valid route or exhausted exact search establishes feasibility status; cases without either certificate remain unresolved.

### B. Dynamics, Costs, and Observation Accounting

Each rest-to-rest SOOA lasts 90 s and uses fixed-step fourth-order Runge–Kutta (RK4) propagation with  $\Delta t = 3$  s. The HCW mean motion is  $1.13137 \times 10^{-3}$  rad/s, the acceleration limit is  $u_{\max} = 0.060$  m/s<sup>2</sup>, and the speed limit is  $v_{\max} = 2.0$  m/s. The vehicle radius is 0.80 m, the full-mesh safety margin is 2 m, and terminal position and velocity tolerances are 0.5 m and 0.05 m/s. Equation (8) retains  $(\lambda_v, \lambda_e, \lambda_s, \lambda_a) = (5, 0.18, 120, 0.05)$ . The action budget is 14. Required-only completion has no additional global area threshold.

Target-identity unions exclude duplicate credit. Reporting includes required-target completion  $C^{\mathcal{K}}$ , priority-weighted inspectable score  $C^{\text{insp}}$ , and unweighted fractions over the 41 observable and 90 original samples. Equal represented sample areas make the last fraction the sampled-area estimate; priority multipliers affect only the weighted score.

### C. Development History and Fresh Confirmation

The first required-target campaign evaluated three fixed profiles on 36 validation, 150 test, and 90 shifted scenarios without feasibility rejection. The nine-item profile was designated primary. Depth-three ADP completed 37/50 test and 10/30 shifted tasks on that profile, whereas greedy construction and its sequence-improvement descendants completed none. Table IV retains these development results, which motivated a completion-capable comparator.

TABLE II

Frozen synthetic required-target profiles. Three equal-width slabs span the mesh longest axis; each slab uses centroid-first, normalized farthest-point selection. IDs omit the common mesh prefix. The nine-item profile is primary

| Profile    | Items | Slab counts | Required sample IDs                                                                |
|------------|-------|-------------|------------------------------------------------------------------------------------|
| required06 | 6     | 2/2/2       | 00055, 00052, 00007, 00001, 00033, 00015                                           |
| required09 | 9     | 3/3/3       | 00055, 00052, 00059, 00007, 00001, 00079, 00033, 00015, 00019                      |
| required12 | 12    | 4/4/4       | 00055, 00052, 00059, 00064, 00007, 00001, 00079, 00045, 00033, 00015, 00019, 00024 |

A fresh confirmation compares depth-three ADP with one-step ADP and local search initialized by a completed one-step ADP route. Methods, penalties, scenario identities, and implementation hashes were frozen before evaluating the new outcomes. The confirmation retains the nine prescribed targets and depth-three policy, with 12 validation, 50 test, and 30 shifted scenarios. Its seeds are disjoint from all development runs, and every generated case is retained.

Each scenario samples a node-removal probability uniformly from  $[0, 0.10]$ , then removes candidates independently using that probability. It multiplies priorities for approximately 6%–14% of targets by a factor in  $[2, 4]$ . Shifted scenarios use removal probabilities in  $[0.15, 0.25]$ , priority fractions of 12%–20%, and factors in  $[5, 8]$ . Priorities affect background scores and the coverage-only ablation, but not the required-target goal or stage cost. All scenarios perturb one frozen graph, and distinct seeds may yield identical routes.

After confirmation, the depth diagnostic was extended to  $d = 1, \dots, 6$  on the same 12 validation scenarios. Every depth used the fixed nine-target requirement, graph, budget, and implementation, including zero-gain connecting actions. The extension retains all cases and leaves the primary test and shifted comparisons unchanged. Fig. 3 shows the resulting cost, runtime, and safety-screen workload.

Complete-route timing includes preliminary base-policy construction and excludes graph loading and serialization. Each scenario–depth pair was measured once in ascending-depth order using Python 3.14 on an Apple M4/macOS host. Scheduling and thermal effects were uncontrolled. Safety-screen counts exclude the separate preliminary base-policy call.

All six depths completed the same nine validation tasks; the other three lacked required visibility. Depth three reduced mean graph cost by 11.19% relative to depth two. It captured 84.11% of the observed reduction from depth one to the lowest mean cost among the six depths. Depth four improved cost by a further 2.74%, at 4.63 times the median planning time and 4.74 times the mean safety-screen count. Depths five and six produced no further cost or completion improvement, while median runtime reached 17.50 and 52.52 times the depth-three value.

The fixed depth-three setting therefore offers a computation–cost compromise; depth four provides lower cost when additional computation is acceptable.




Fig. 3. Post-selection depth diagnostic. The square marks the primary depth-three setting. Cost uses the nine tasks completed at all depths; runtime and safety-screen counts include all 12 validation cases. Panels (b) and (c) use logarithmic scales. Table III reports completion counts and changes relative to depth three.

## D. Matched Baselines and Statistical Reporting

The confirmation includes the task-aware greedy incumbent, one-step ADP, ADP-seeded local search, and depth-three ADP under identical tasks, dynamics, safety gates, and action budgets. One-step ADP uses (17) with  $d = 1$ . Local search applies three audited best-improvement reversal passes to its completed route and retains the feasible seed if improvement fails. This comparator improves an ADP-initialized ordering; it does not introduce new viewpoints. Depth-three ADP uses the original greedy base value without a one-step fallback.

A coverage-only ablation retains the earlier 80% priority-weighted stopping rule and is evaluated against the same required-target identities. It tests the mission objective separately from the matched complete-task cost comparison.

Success and missing-target counts use every scenario; graph cost and physical  $\Delta v$  use jointly completed ADP/seeded-local pairs. The all-scenario score is  $J_{\text{pen}} = J + 500(1 - C^K) + 500(1 - \Gamma_K)$ , with fixed evaluation penalties distinct from physical maneuver cost. Paired-mean percentile intervals use 10 000 resamples, and exact two-sided sign tests exclude numerical ties. Timing includes initialization, including the one-step seed for local search; safety-screen counts cover the selected planner call.

## VI. Results

### A. Complete Required-Target Inspection

Depth-three ADP completed all nine targets in 43/50 test scenarios and 9/30 shifted scenarios. Seeded local

TABLE III

Post-selection required-target depth diagnostic on the same 12 validation scenarios. Graph-cost means use the 9 scenarios completed at all six depths. Time and screen counts include all 12 cases; ratios use depth three. Negative  $\Delta J$  denotes lower mean cost. All depths were rerun sequentially; the primary depth-three test results are unchanged

| Depth    | Complete    | Mean $J$       | $\Delta J$ (%) | Median time (s) | Time ( $\times$ ) | Mean screens   | Screens ( $\times$ ) |
|----------|-------------|----------------|----------------|-----------------|-------------------|----------------|----------------------|
| 1        | 9/12        | 135.714        | +14.51         | 0.0272          | 0.03              | 7,132          | 0.02                 |
| 2        | 9/12        | 133.450        | +12.60         | 0.2207          | 0.22              | 65,928         | 0.20                 |
| <b>3</b> | <b>9/12</b> | <b>118.520</b> | <b>+0.00</b>   | <b>1.0094</b>   | <b>1.00</b>       | <b>322,349</b> | <b>1.00</b>          |
| 4        | 9/12        | 115.272        | -2.74          | 4.6760          | 4.63              | 1,527,408      | 4.74                 |
| 5        | 9/12        | 115.272        | -2.74          | 17.6601         | 17.50             | 5,757,081      | 17.86                |
| 6        | 9/12        | 115.272        | -2.74          | 53.0133         | 52.52             | 17,454,611     | 54.15                |

search completed 43/50 and 8/30, respectively. Every completed mission achieved  $C^{\mathcal{K}} = 100\%$ , with optional background coverage reported separately.

ADP completed every confirmation scenario retaining visibility of all required items. Each remaining case lost all available views of at least one mandatory target. In the development campaign, one six-target shifted case retained required visibility but produced no ADP certificate; its directed feasibility remains unresolved.

#### B. ADP Policy Improvement at Matched Completion

The predeclared ADP-minus-seeded-local penalized endpoint had mean differences of  $-12.256$  on the 50 test scenarios and  $-31.999$  on the 30 shifted scenarios. Their 95% intervals were  $[-14.155, -10.212]$  and  $[-91.494, -1.241]$ , respectively. Physical-resource comparisons below use the jointly completed missions.

Fig. 4 compares the 43 jointly completed test pairs and 8 shifted pairs. Mean ADP-minus-local graph-cost differences were  $-14.251$  and  $-9.890$ , respectively. Table VI reports the paired uncertainty and win/tie/loss counts.

On the 43 completed test pairs, mean  $\Delta v$  decreased from  $15.777$  m/s to  $14.033$  m/s, an 11.05% reduction. The paired reduction was  $1.744$  m/s, with a 95% interval of  $1.525$ – $1.923$  m/s; ADP won 40 pairs and lost three. The eight shifted pairs had a mean reduction of  $1.207$  m/s. Its bootstrap interval excluded zero, while the exact sign test gave  $p = 0.0703$ .

The greedy base completed no confirmation mission, so this campaign measures improvement over one-step ADP and seeded local search rather than a completing greedy comparator. Finite-graph unit tests separately exercise the base-cost bound, connecting actions, and exact recursion.

#### C. Objective Ablation and Safety

Fig. 5 summarizes component performance and the accepted-edge safety metrics. The coverage-only ablation in Table VII applies the fixed required-target denominator to routes using the 80% background stopping rule.

The coverage-only policy completed the nine-item requirement in 1/50 test and 0/30 shifted scenarios. In contrast, 41 of 43 completed ADP test missions finished

below 80% weighted background coverage while accepting all nine prescribed items. Mean background coverage over those 43 successes was 72.32%.

Full-mesh occlusion and collision checks include structure outside the required target set. The reported clearance and input metrics describe the accepted deterministic edge records.

#### D. Representative Required-Target Route

Fig. 6 shows the selected routes, target completion, and accumulated  $\Delta v$  for the representative test pair.

The illustrative pair is scenario `required09_test_000`, selected as the jointly complete test pair nearest the median ADP-minus-local graph-cost difference (deterministic ID tie break). Both routes complete all 9 prescribed targets. Rollout ADP uses 9 SOOAs,  $J = 113.472$ , and  $\Delta v = 13.769$  m/s; local search uses 8 SOOAs,  $J = 129.605$ , and  $\Delta v = 15.744$  m/s. The respective background inspectable scores are 77.22% and 64.97%. Repropagation reproduces both archived costs and velocity increments to numerical tolerance.

#### E. Execution Interface and Historical ROS Evidence

The execution interface records per-observation target IDs and unions accepted observations. Tests verify required completion, duplicate exclusion, rejection handling, and export consistency. The schema preserves existing message interfaces while distinguishing planned and accepted metrics.

An offline replay export preserves the illustrative route’s nine actions, nine required IDs, and 77.22% planned background coverage. Target unions and resource metrics match the archive. Six CSV streams record trajectory, control, coverage, safety, planning, and mission events.

The available Ubuntu 24.04.4/ROS 2 Jazzy execution uses the earlier 80% survey objective and is reported separately below. Required-target results in this study remain offline planning and interface tests; a new required-target ROS campaign has not been executed.

TABLE IV

First-campaign development results by fixed profile. Local search in this campaign was initialized by the greedy incumbent, unlike the fresh confirmation comparator. Profiles use different seeds, so cross-profile rates are descriptive

| Profile    | Required items | Split   | ADP complete | Local complete |
|------------|----------------|---------|--------------|----------------|
| required06 | 6              | Test    | 40/50        | 0/50           |
| required09 | 9              | Test    | 37/50        | 0/50           |
| required12 | 12             | Test    | 38/50        | 0/50           |
| required06 | 6              | Shifted | 10/30        | 0/30           |
| required09 | 9              | Shifted | 10/30        | 0/30           |
| required12 | 12             | Shifted | 10/30        | 1/30           |

TABLE V

Required-target results. Resource means use successful routes of each method; the final column includes every scenario with the fixed failure penalty

| Split   | Method                  | Complete | Mean SOOAs | Mean $\Delta v$ (m/s) | Mean $J$ | Penalized $J$ |
|---------|-------------------------|----------|------------|-----------------------|----------|---------------|
| Test    | Task-aware incumbent    | 0/50     | —          | —                     | —        | 911.242       |
| Test    | One-step ADP            | 43/50    | 8.21       | 16.512                | 135.945  | 256.913       |
| Test    | ADP-seeded local search | 43/50    | 8.21       | 15.777                | 129.892  | 251.707       |
| Test    | Depth-three ADP         | 43/50    | 9.00       | 14.033                | 115.641  | 239.451       |
| Shifted | Task-aware incumbent    | 0/30     | —          | —                     | —        | 851.059       |
| Shifted | One-step ADP            | 8/30     | 9.25       | 17.012                | 140.102  | 770.694       |
| Shifted | ADP-seeded local search | 8/30     | 8.88       | 15.888                | 130.841  | 768.224       |
| Shifted | Depth-three ADP         | 9/30     | 9.11       | 14.656                | 120.751  | 736.225       |

TABLE VI

ADP-minus-seeded-local effects. The predeclared penalized endpoint includes all scenarios; physical resources and graph cost use jointly complete pairs. Intervals are percentile 95% paired-bootstrap intervals; sign tests exclude ties

| Split   | Endpoint               | Pairs | Mean difference | 95% interval         | Win/tie/loss | Sign-test $p$         |
|---------|------------------------|-------|-----------------|----------------------|--------------|-----------------------|
| Test    | $J_{\text{pen}}$ (all) | 50    | -12.256         | $[-14.155, -10.212]$ | 40/7/3       | $3.02 \times 10^{-9}$ |
| Test    | $J$                    | 43    | -14.251         | $[-15.715, -12.461]$ | 40/0/3       | $3.02 \times 10^{-9}$ |
| Test    | $\Delta v$ (m/s)       | 43    | -1.744          | $[-1.923, -1.525]$   | 40/0/3       | $3.02 \times 10^{-9}$ |
| Shifted | $J_{\text{pen}}$ (all) | 30    | -31.999         | $[-91.494, -1.241]$  | 8/21/1       | 0.0391                |
| Shifted | $J$                    | 8     | -9.890          | $[-14.111, -4.688]$  | 7/0/1        | 0.0703                |
| Shifted | $\Delta v$ (m/s)       | 8     | -1.207          | $[-1.713, -0.575]$   | 7/0/1        | 0.0703                |





Fig. 4. ADP/seeded-local comparison on jointly completed missions: 43 test and 8 shifted pairs. Dashed lines denote parity. Bands and horizontal lines in (b) and (d) show 95% paired-bootstrap intervals and mean graph-cost differences.


## F. Historical Corrected ROS 2 Survey Execution

The historical 80%-survey validation task ran on Ubuntu 24.04.4 with ROS 2 Jazzy, Gazebo Harmonic, and RViz2. Gazebo supplied the station/chaser rendering and camera image only; the ROS-native HCW node remained the state source of truth. The input route was reconstructed

from the corrected 24-node graph using the complete gITF scene hierarchy and a conservative  $r_{\text{veh}} = 0.80$  m chaser bounding radius. The route manifest records 10 planned SOOAs, 900 s mission duration, 81.33% planned coverage, and 2.808 m minimum planned finite-body clearance above the required 2 m margin.



(a) Matched complete-task effort. (b) All-required mission success.




(c) Clearance above margin. (d) Peak commanded input.

Fig. 5. Nine-target component and safety comparison. Panel (a) shows means and 95% bootstrap intervals for the 43 test cases completed by one-step ADP, seeded local search, and depth-three ADP. Panel (b) retains all test and shifted cases. Panels (c) and (d) use jointly completed ADP/seeded-local pairs.


TABLE VII

Required-target completion under the coverage-only objective. The required set is identical to that of the main comparison

| Split   | Objective ablation | All required complete |
|---------|--------------------|-----------------------|
| Test    | coverage only80    | 1/50                  |
| Shifted | coverage only80    | 0/30                  |



(a) Generated LVLH routes.



(b) Required-target completion. (c) Cumulative physical  $\Delta v$ .

Fig. 6. Representative required-target test pair selected by the median-effect rule. Both methods complete the same required set. Markers identify terminal observations and prescribed samples; the star denotes the common initial state.

The  $5 \times$  graphical closed-loop execution credited all 10 terminal observations, reached 81.33% coverage, and completed at mission time 900.01 s without a failed action. The post-run audit evaluated 18918 recorded trajectory samples and every intervening segment against all 247525 fully transformed ISS triangles. No segment crossed the mesh. The smallest sampled center distance was 8.506 m; after the continuous segment lower bound and the 0.80 m vehicle radius, the minimum body distance was 7.692 m, or 5.692 m above the required safety mar-




Fig. 7. Historical survey execution: trajectory-linked camera views from the accepted corrected validation\_002 task with an 80% inspectable-coverage goal. Views 1–5 link locally to the radial-cross-track ( $x$ – $z$ ) projection, and views 6–10 link to the along-track-cross-track ( $y$ – $z$ ) projection; hollow gray markers retain the other observations as context. Inset labels give mission time and cumulative coverage. Dashed gray and solid blue denote planned and executed paths, and the star is the initial state. Clearance is evaluated against the full 3-D mesh and swept path. Gazebo renders the views, ROS HCW dynamics propagate state, and frozen geometric visibility masks determine credited coverage.

gin. Peak filtered acceleration was  $0.0600 \text{ m/s}^2$ . Terminal, reference-stream, data-presence, acceleration, full-mesh, swept-crossing, and finite-body gates all passed.

Fig. 7 links the ten synchronized camera frames to their credited positions in two orthogonal projections of the executed trajectory. The evidence directory contains six CSV streams, configuration and route manifests, the launch log, mesh audit, raw RViz and camera streams, and a 198.3-s annotated video.

Earlier translation-only ROS results are excluded because the complete glTF hierarchy exposed an S4-truss intersection and negative finite-body clearance. The retained execution uses the corrected geometry and audit.

## VII. Discussion

The results support model-based rollout ADP for prescribed-target inspection. A base-policy completion supplies a checkable continuation, while Bellman lookahead improves the next action. The required-target con-

straint determines mission completion independently of optional background coverage.

The task formulation also accommodates a hybrid goal with an additional background threshold, provided the base policy and rollout boundary use the same predicate. Target priorities alone cannot impose mandatory observations. The coverage-only ablation illustrates this difference directly.

The empirical comparison measures maneuver savings against one-step ADP and its seeded reversal-search baseline. Equal-computation comparisons with broader search neighborhoods, exact optimality gaps, and larger independent graphs would establish the method's comparative range more fully.

## A. Scope and Limitations

The evidence covers synthetic prescribed samples on one deterministic observation graph. Completion certifies those samples under the camera model, rather than full-region imagery or defect detection. Operational inspection requires verified component requirements, denser sampling, resolution and dwell acceptance, illumination, and independent geometry.

The finite-completion safeguard may fail to find a route when its base policy is unsuccessful, even if another feasible route exists. The base-cost theorem applies when the base policy completes; exact optimality applies only to an exhausted graph. Timing and confidence intervals characterize the stated desktop implementation and scenario generator. The observed depth-four-to-six plateau does not establish global optimality.

The model assumptions in Table I exclude navigation uncertainty, actuator and attitude dynamics, and illumination effects. Safety remains conditional on the enabled deterministic audits, with passive-drift auditing disabled. The study evaluates model-based ADP; extending it to a learned critic or actor would require preserving a valid continuation and the relevant cost bound.

## VIII. Conclusion

This paper developed viability-preserving rollout ADP for prescribed-target inspection over audited orbital arcs. A finite Bellman prefix and task-aware base-policy value preserve completion and provide a conditional graph-cost bound.

ADP completed 43/50 test and 9/30 shifted missions. On jointly completed test missions, it reduced mean  $\Delta v$  by 11.05% relative to seeded local search. The depth study quantified the additional computation required for further cost reduction. These results support completion-certified sequencing within the prescribed finite inspection model.

## Data and Code Availability

The accompanying reproducibility package contains the corrected graph, fixed target profiles and their ge-

ometric provenance, configuration and freeze manifests, fresh scenario inventories, method-scenario rows, target-level credit records, statistical outputs, source tables, and figure scripts. Required-target evidence is versioned separately from the historical 80% benchmark, post-selection depth/coverage diagnostics, and accepted historical ROS task. A new required-target ROS bag is not included because that campaign was not run. The public NASA mesh source and checksum are documented; the mesh is not redistributed. A persistent repository identifier will replace the submission-package reference in the accepted version.

## Acknowledgment

OpenAI Codex was used to assist with drafting and revision in the Abstract, Introduction, System and Observation Model, Viability-Preserving Rollout ADP, Experimental Protocol, Results, Discussion, Conclusion, and Data and Code Availability sections; with LaTeX formatting and source checking; and with implementation, testing, offline experiment orchestration, statistical tabulation, and figure-generation scripts. The authors retain responsibility for reviewing and validating all AI-assisted text, mathematical statements, references, source code, figures, and reported results, and for the scientific content of the article.

## REFERENCES

- [1] S. Faghihi, S. Tavana, and A. H. J. de Ruiter, "Kinodynamic on-orbit inspection path planning for full-coverage inspection in close proximity of space structures," *Acta Astronautica*, vol. 198, pp. 354–365, 2022.
- [2] —, "Multiple spacecraft coordination and motion planning for full-coverage inspection of large complex space structures," *Acta Astronautica*, vol. 202, pp. 119–129, 2023.
- [3] B. Apodaca, T. Helgeson, E. Atkins, and L. Stirling, "In-orbit space structure inspection trajectory generation," *IEEE Robotics and Automation Letters*, vol. 10, no. 9, pp. 8810–8817, 2025.
- [4] W. R. Scott, G. Roth, and J.-F. Rivest, "View planning for automated three-dimensional object reconstruction and inspection," *ACM Computing Surveys*, vol. 35, no. 1, pp. 64–96, 2003.
- [5] Z. Hou, B. Jiao, and Z. Dang, "Relative orbit design of cubesats for on-orbit visual inspection of china space station," *Advances in Space Research*, vol. 73, no. 1, pp. 170–186, 2024.
- [6] T. Deka and J. W. McMahon, "Astrodynamics-informed kinodynamic sampling-based motion planning for relative spacecraft motion," *Journal of Guidance, Control, and Dynamics*, vol. 46, no. 12, pp. 2330–2345, 2023.
- [7] D. P. Bertsekas, *Rollout, Policy Iteration, and Distributed Reinforcement Learning*. Athena Scientific, 2020.
- [8] W. H. Clohessy and R. S. Wiltshire, "Terminal guidance system for satellite rendezvous," *Journal of the Aerospace Sciences*, vol. 27, no. 9, pp. 653–658, 1960.
- [9] J. R. Wertz, D. F. Everett, and J. J. Puschell, *Space Mission Engineering: The New SMAD*. Microcosm Press, 2011.
- [10] H. Zhou, B. Jiao, Z. Hou, and Z. Dang, "Satellite formation control for cooperative observation around the rotating china space station," *Space: Science & Technology*, vol. 6, p. 0348, 2026.
- [11] S. M. LaValle, "Rapidly-exploring random trees: A new tool for path planning," Department of Computer Science, Iowa State University, Tech. Rep. TR 98-11, October 1998.

- [12] S. Karaman and E. Frazzoli, “Sampling-based algorithms for optimal motion planning,” *The International Journal of Robotics Research*, vol. 30, no. 7, pp. 846–894, 2011.
- [13] D. P. Bertsekas and J. N. Tsitsiklis, *Neuro-Dynamic Programming*. Athena Scientific, 1996.
- [14] J. Achiam, D. Held, A. Tamar, and P. Abbeel, “Constrained policy optimization,” in *Proceedings of the 34th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 70, 2017, pp. 22–31.
- [15] M. Alshiekh, R. Bloem, R. Ehlers, B. Könighofer, S. Niekum, and U. Topcu, “Safe reinforcement learning via shielding,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, 2018, pp. 2669–2678.
- [16] J. Zhou, G. Kang, J. Wu, X. Tao, C. Xu, Y. Qiu, and J. Wu, “Tomographic inspection path planning for spacecraft with complex structures,” *Acta Astronautica*, vol. 239, pp. 162–183, 2026.
- [17] T. E. Ogri, M. Qureshi, Z. I. Bell, W. A. Makumi, K. Volle, and R. Kamalapurkar, “A switched systems approach to image-based feature tracking for autonomous satellite inspection,” 2025.
- [18] National Aeronautics and Space Administration, “International Space Station glTF Model,” <https://science.nasa.gov/resource/international-space-station-3d-model/>, 2019, updated Nov. 2, 2024; accessed Aug. 10, 2026.